

diabetes

September 24, 2024

```
[1]: import pandas as pd
import numpy as np
```

1 Importation des données à partir d'un fichier csv présent dans le dossier

```
[2]: data = pd.read_csv("diabetes.csv")
data.head()
```

```
[2]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

1.1 Identification des caractéristiques et de la variable cible

```
[3]: features_names = ["Pregnancies", "Glucose", "BloodPressure", "SkinThickness", "Insulin", "BMI", "DiabetesPedigreeFunction", "Age"]
X = data[features_names]
y = data["Outcome"]
```

2 Exploration

```
[4]: X.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Pregnancies           768 non-null   int64
1   Glucose               768 non-null   int64
2   BloodPressure         768 non-null   int64
3   SkinThickness         768 non-null   int64
4   Insulin               768 non-null   int64
5   BMI                   768 non-null   float64
6   DiabetesPedigreeFunction 768 non-null   float64
7   Age                   768 non-null   int64
dtypes: float64(2), int64(6)
memory usage: 48.1 KB

```

Il n'y a pas de valeur manquante, tous les types des données sont numériques.

```
[5]: X.describe()
```

```

[5]:      Pregnancies    Glucose  BloodPressure  SkinThickness    Insulin  \
count    768.000000   768.000000    768.000000    768.000000   768.000000
mean      3.845052   120.894531     69.105469     20.536458    79.799479
std       3.369578    31.972618     19.355807     15.952218   115.244002
min       0.000000     0.000000     0.000000     0.000000     0.000000
25%       1.000000    99.000000     62.000000     0.000000     0.000000
50%       3.000000   117.000000     72.000000     23.000000    30.500000
75%       6.000000   140.250000     80.000000     32.000000   127.250000
max      17.000000   199.000000    122.000000     99.000000   846.000000

      BMI  DiabetesPedigreeFunction    Age
count    768.000000          768.000000   768.000000
mean     31.992578              0.471876    33.240885
std       7.884160              0.331329    11.760232
min       0.000000              0.078000    21.000000
25%      27.300000              0.243750    24.000000
50%      32.000000              0.372500    29.000000
75%      36.600000              0.626250    41.000000
max      67.100000              2.420000    81.000000

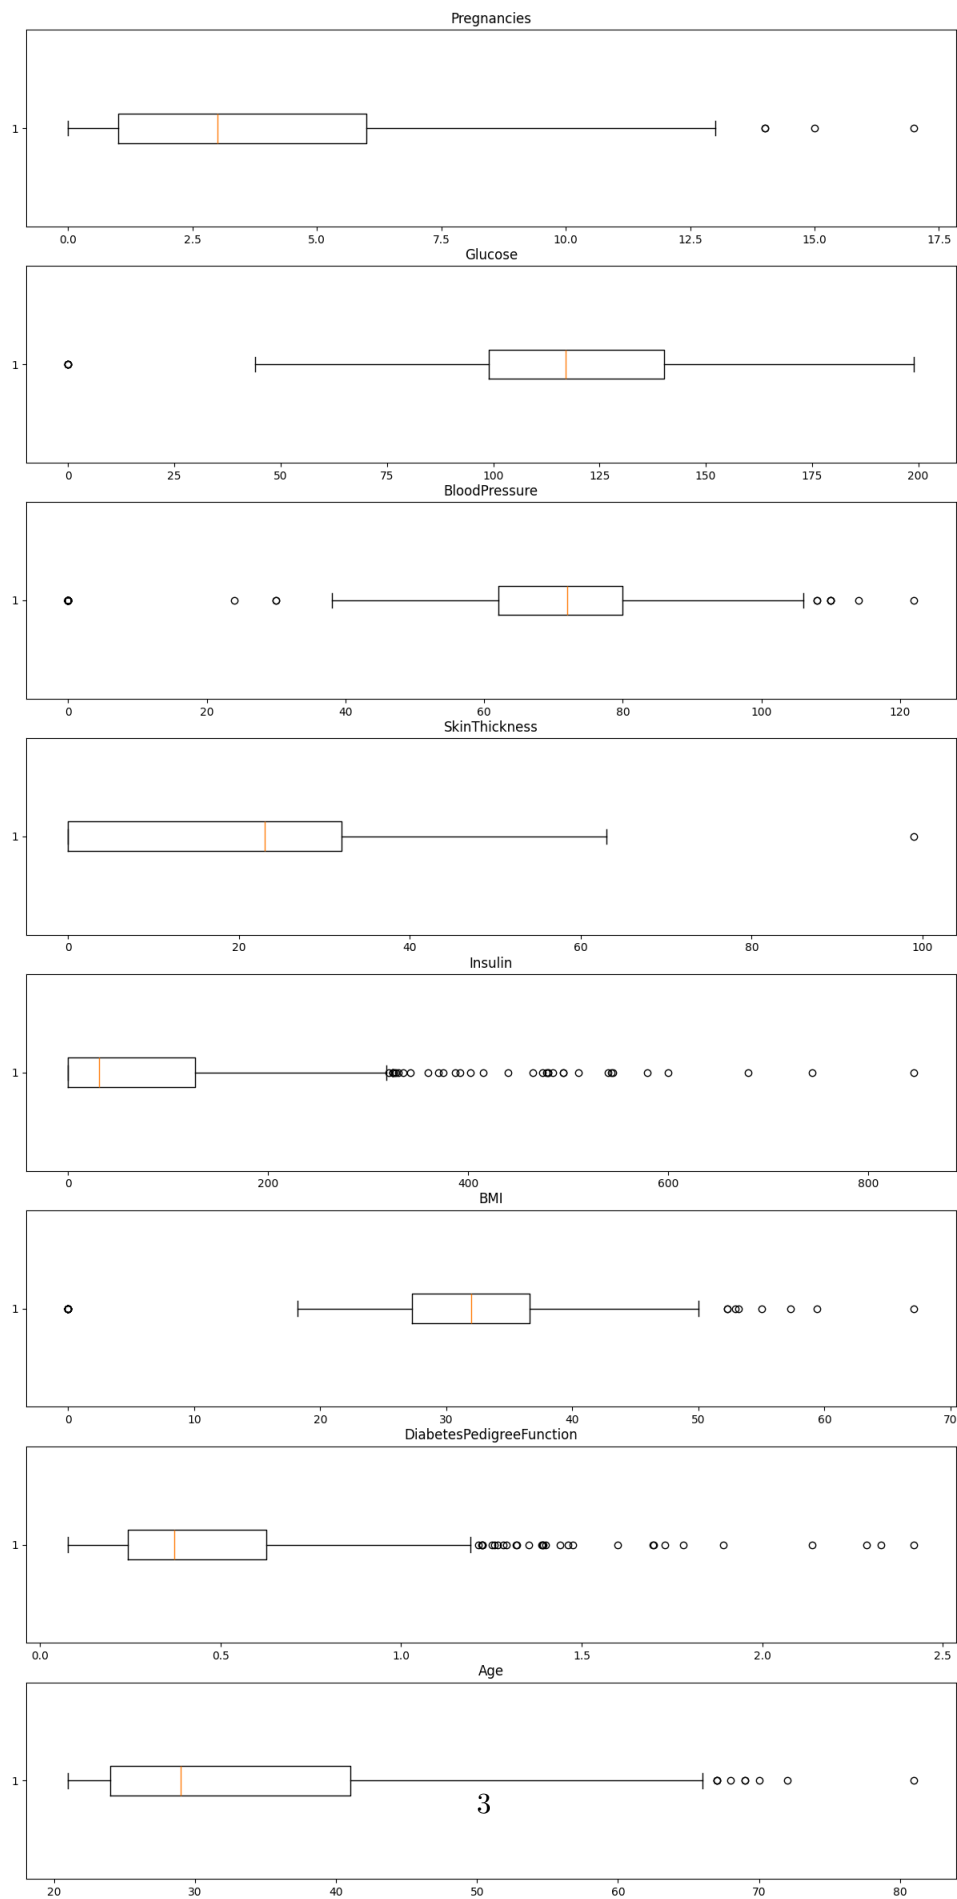
```

```

[6]: import matplotlib.pyplot as plt

plt.figure(figsize=(15,30))
for i, name in enumerate(features_names):
    plt.subplot(len(features_names), 1, i+1)
    plt.title(name)
    plt.boxplot(data[name], vert=False)

```



```
[7]: data.loc[data["Outcome"] == 1, "Outcome"].count()
```

```
[7]: 268
```

```
[8]: data.loc[data["Outcome"] == 0, "Outcome"].count()
```

```
[8]: 500
```

Il y a deux fois moins de données pour les cas de diabète par rapport aux cas négatifs ce qui peut biaiser le résultat. La solution est d'enrichir la classe minoritaire ou d'appauvrir la majoritaire. Ici on a choisi le 1er cas pour ne pas perdre de données.

2.0.1 Séparation en jeu d'entraînement et de test et création de nouveaux exemples synthétiques pour la classe minoritaire (ici outcome=1 donc cas diabétiques)

```
[9]: from sklearn.model_selection import train_test_split
from imblearn.over_sampling import SMOTE

X_train, X_test, y_train, y_test = train_test_split(data[features_names],
    ↪data["Outcome"], test_size=0.2, random_state=24 )
smote = SMOTE()
X_resampled, y_resampled = smote.fit_resample(X_train, y_train)
```

2.0.2 Pour que la classification se passe bien nous allons normaliser et standardiser les données.

```
[10]: from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
X_resampled_scaled = scaler.fit_transform(X_resampled)
X_test_scaled = scaler.transform(X_test)
```

3 Utilisant la validation croisée puis entraînement du modèle

3.1 1 Modèle de régression logistique

```
[11]: from sklearn.model_selection import cross_val_score
from sklearn.linear_model import LogisticRegression

model = LogisticRegression()
scores = cross_val_score(model, X_resampled_scaled, y_resampled, cv=5)
print(f'Mean accuracy: {scores.mean()}')
print(f'Accuracy: {scores}')
```

```
model.fit(X_resampled_scaled, y_resampled)
```

Mean accuracy: 0.7474534161490685

Accuracy: [0.74534161 0.82608696 0.69565217 0.77018634 0.7]

```
[11]: LogisticRegression()
```

```
[12]: y_pred = model.predict(X_test_scaled)
```

3.1.1 Affichage des différents scores

```
[21]: from sklearn.metrics import precision_score, recall_score, f1_score, \
      ↪ confusion_matrix, classification_report

def print_metric(y_test, y_pred):
    precision = precision_score(y_test, y_pred)
    recall = recall_score(y_test, y_pred)
    f1 = f1_score(y_test, y_pred)
    cm = confusion_matrix(y_test, y_pred)

    print(f"Precision score: {precision}")
    print(f"Recall score: {recall}")
    print(f"f1 score: {f1}")
    print(f"Confusion matrix: \n{cm}")

    print("Rapport de classification :\n", classification_report(y_test, \
    ↪ y_pred))

print_metric(y_test, y_pred)
```

Precision score: 0.6481481481481481

Recall score: 0.625

f1 score: 0.6363636363636364

Confusion matrix:

```
[[79 19]
```

```
[21 35]]
```

Rapport de classification :

	precision	recall	f1-score	support
0	0.79	0.81	0.80	98
1	0.65	0.62	0.64	56
accuracy			0.74	154
macro avg	0.72	0.72	0.72	154
weighted avg	0.74	0.74	0.74	154

3.2 2 Algorithmes des K-Nearest Neighbors (KNN)

```
[18]: from sklearn.neighbors import KNeighborsClassifier
      from sklearn.model_selection import GridSearchCV

      knn = KNeighborsClassifier()
      param_grid = {'n_neighbors': [2, 3, 4, 5, 7, 9]}
      grid_search = GridSearchCV(knn, param_grid, cv=5)
      grid_search.fit(X_resampled_scaled, y_resampled)
      print(f"Best k: {grid_search.best_params_}")
```

Best k: {'n_neighbors': 5}

```
[19]: knn = KNeighborsClassifier(n_neighbors=5)
      knn.fit(X_resampled_scaled, y_resampled)
      y_knn_pred = knn.predict(X_test_scaled)
```

3.2.1 Affichage des différents scores

```
[22]: print_metric(y_test, y_knn_pred)
```

Precision score: 0.6229508196721312

Recall score: 0.6785714285714286

f1 score: 0.6495726495726496

Confusion matrix:

```
[[75 23]
 [18 38]]
```

Rapport de classification :

	precision	recall	f1-score	support
0	0.81	0.77	0.79	98
1	0.62	0.68	0.65	56
accuracy			0.73	154
macro avg	0.71	0.72	0.72	154
weighted avg	0.74	0.73	0.74	154

L'algorithme des k plus proches voisins a plus de précision mais un moins bon rappel, il y a moins de faux positifs mais plus de faux négatifs.

```
[43]: from sklearn.ensemble import RandomForestClassifier

      param_grid = {
          'n_estimators': [50, 100, 150, 200, 300],
          'max_depth': [2, 40, 50, 70, 100, 150],
          'max_features': [None, 'sqrt', 'log2']
      }
```

```
grid_search = GridSearchCV(RandomForestClassifier(random_state=24), param_grid,
    ↪cv=5, scoring='recall')
grid_search.fit(X_resampled_scaled, y_resampled)
print(f'Best parameters: {grid_search.best_params_}')
```

Best parameters: {'max_depth': 40, 'max_features': 'sqrt', 'n_estimators': 300}

```
[39]: rf = RandomForestClassifier(n_estimators=300, random_state=24,
    ↪max_features='sqrt', max_depth=40)
rf.fit(X_resampled_scaled, y_resampled)
y_rf_pred = rf.predict(X_test_scaled)
```

```
[40]: print_metric(y_test, y_rf_pred)
```

Precision score: 0.5849056603773585

Recall score: 0.5535714285714286

f1 score: 0.5688073394495413

Confusion matrix:

```
[[76 22]
 [25 31]]
```

Rapport de classification :

	precision	recall	f1-score	support
0	0.75	0.78	0.76	98
1	0.58	0.55	0.57	56
accuracy			0.69	154
macro avg	0.67	0.66	0.67	154
weighted avg	0.69	0.69	0.69	154

Le modèle des arbres a aussi plus de précision mais un moins bon rappel, il y a moins de faux positifs mais plus de faux négatifs.

```
[ ]:
```